

Jack's car rental

Corrado Possieri

Machine and Reinforcement Learning in Control Applications

Problem



Manage cars between two locations.

Problem formulation

- Jack manages two locations for a car rental company.
- Each day, some customers arrive at each location to rent cars.
- If Jack has a car available, he rents it out and is credited \$10.
- Cars are available for renting the day after they are returned.
- Jack can move cars between the two locations overnight.
- The cost of moving a car is \$2.
- Each location is capable of accommodating 30 cars.





Problem data

- Cars requested and returned at each location are Poisson random variables

$$\mathbb{P}[\text{cars} = n] = \frac{\lambda^n}{n!} \exp(-\lambda).$$

- λ is 3 and 4 for rental requests.
- λ is 3 and 2 for returns.
- Jack's foresight modeled with discount $\gamma = 0.9$.
- Jack can move up to 7 cars between the two locations.

Model

- We can model the process as an MDP.
- The state is the number of car at each location
 - is it a Markov state?
 - we have $\#1 \cdot \#2$ states. 
- The action is the number of cars moved
 - we have $2\#c + 1$ actions. 

State update

- ① Jack reintroduces car returned at previous day.
- ② Jack rents available cars.
- ③ Jack moves cars between the two locations.



Transition and reward

- Let S and A be the number of states and actions.
- Transition probabilities can be stored in a $S \times S \times A$ matrix P

$$P_{s,s',a} = \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a]$$

- Rewards can be stored in a $S \times A$ matrix R

$$R_{s,a} = \mathbb{E}[R_{t+1} | S_t = s, A_t = a]$$

Transition probabilities

- Probability of returns do not depend on actions
 - define the return probability matrix P_{ret}

$$[P_{\text{ret}}]_{s,s'} = \mathbb{P}[S_{\text{after return}} = s' | S_t = s].$$

- Probability of rentals do not depend on actions
 - define the rental probability matrix P_{ren}



$$[P_{\text{ren}}]_{s,s'} = \mathbb{P}[S_{\text{after rental}} = s' | S_{\text{after return}} = s].$$

- Probability of movement depend on actions
 - define the movement probability matrix P_{mov}

$$[P_{\text{mov}}]_{s,s',a} = \mathbb{P}[S_{\text{after movement}} = s' | S_{\text{after rental}} = s, A_t = a].$$

- By the law of total probability

$$P = P_{\text{ret}} \cdot P_{\text{ren}} \cdot P_{\text{mov}}.$$

Expected rewards

- Expected earning do not depend on action

$$\mathbb{E} [\text{earning}_{t+1} | S_{\text{after return}} = s] = \sum_r r \mathbb{P} [r | S_{\text{after return}} = s] .$$

- By the law of total probability

$$\begin{aligned} & \mathbb{E} [\text{earning}_{t+1} | S_t = s] \\ &= \sum_{s'} \mathbb{P} [S_{\text{after return}} = s' | S_t = s] \sum_r r \mathbb{P} [r | S_{\text{after return}} = s'] . \end{aligned}$$

- The expected reward is given by

$$R_{s,a} = [P_{\text{ret}} \cdot \text{earning}]_s - \text{cost}_a .$$

PI and VI revisited

- Given a deterministic policy π , define

$$\begin{aligned}
 \blacksquare P^\pi &= \begin{bmatrix} P_{1,1,\pi(1)} & P_{1,2,\pi(1)} & \cdots & P_{1,S,\pi(1)} \\ P_{2,1,\pi(2)} & P_{2,2,\pi(2)} & \cdots & P_{2,S,\pi(2)} \\ \vdots & \vdots & \ddots & \vdots \\ P_{S,1,\pi(S)} & P_{S,2,\pi(S)} & \cdots & P_{S,S,\pi(S)} \end{bmatrix} \\
 \blacksquare R^\pi &= \begin{bmatrix} R_{1,\pi(1)} \\ R_{2,\pi(2)} \\ \vdots \\ R_{S,\pi(S)} \end{bmatrix}
 \end{aligned}$$

- Bellman expectation update can be rewritten as


$$v \leftarrow R^\pi + \gamma P^\pi v \quad (v^\pi = (I - \gamma P^\pi) R^\pi).$$

- Bellman optimality update can be rewritten as

$$v \leftarrow \max_{\pi} \{R^\pi + \gamma P^\pi v\}$$

Matrix Formulation

- Recall classical Bellman update


$$\begin{aligned}
 v(s) &\leftarrow \sum_a \pi(a|s) \sum_{s', r} p(s', r|s, a) (r + \gamma v(s')) \\
 &= \sum_{s', r} p(s', r|s, \pi(s)) (r + \gamma v(s')) \\
 &= \sum_{s', r} p(s', r|s, \pi(s)) r + \gamma \sum_{s', r} p(s', r|s, \pi(s)) v(s') \\
 &= \sum_r r \sum_{s'} (s', r|s, \pi(s)) + \gamma \sum_{s'} v(s') \sum_r p(s', r|s, \pi(s)) \\
 &= r(s, \pi(s)) + \gamma \sum_{s'} p(s'|s, \pi(s)) v(s').
 \end{aligned}$$


- A similar relation holds for Bellman optimality update

$$v(s) \leftarrow \max_a \left\{ r(s, a) + \gamma \sum_{s'} p(s'|s, a) v(s') \right\}.$$

Assignment #2

- Write a code for PI (in class).
- Write a code for VI (in class).



- Model tic-tac-toe as an MDP

- each board configuration is a state;
- actions is where to place your mark;
- the other player is playing at random;
- use **afterstates**
 - ▶ evaluate board positions after the other p
- reward +1 for winning and -1 for losing;
- see Section 1.5 of textbook for some hints
- use VI and PI to determine optimal actions.

