

# Markov Decision Processes

Corrado Possieri

Machine and Reinforcement Learning in Control Applications

# Introduction

- **Markov decision processes** (MDP) formally describe an environment for reinforcement learning.
- The environment is fully observable
  - the current state completely characterizes the future.
- Almost all learning problems can be formalized as MDPs:
  - optimal control deals with continuous MDPs;
  - partially observable problems can be converted into MDPs;
  - bandits are MDPs with just a single state.

# Markov property

## Markov property

The future is independent of the past given the present.

Formally

$$\mathbb{P}[S_{t+1}|S_t] = \mathbb{P}[S_{t+1}|S_1, S_2, \dots, S_t].$$

- The state captures all relevant information from the history.
- Once the state is known, the history may be thrown away
  - the state is a sufficient statistic of the future.

# State transition matrix

- Given states  $s$  and  $s'$ , the state transition probability is

$$P_{s,s'} = \mathbb{P}[S_{t+1} = s' | S_t = s].$$

- If the states are **finite**
  - define the state transition matrix

$$P = \begin{matrix} & \text{to} \\ \text{from} & \begin{bmatrix} P_{1,1} & P_{1,2} & \cdots & P_{1,n} \\ \vdots & \vdots & \ddots & \vdots \\ P_{n,1} & P_{n,2} & \cdots & P_{n,n} \end{bmatrix} \end{matrix} ;$$

- each row of  $P$  sums to 1:

$$\sum_{j=1}^n P_{i,j} = 1, \quad i = 1, \dots, n.$$

# Probability distributions in Markov chains

- Letting

$$\pi(t) = \begin{bmatrix} \mathbb{P}[S_t = 1] \\ \mathbb{P}[S_t = 2] \\ \vdots \\ \mathbb{P}[S_t = n] \end{bmatrix}^\top,$$

one has

$$\pi(t+1) = \pi(t) P.$$

- **Stationary distributions** satisfy

$$\bar{\pi} = \bar{\pi} P.$$

- We have that

$$\mathbb{P}[S_{t+h} = s' | S_t = s] = [P^h]_{s',s}.$$

# Markov process

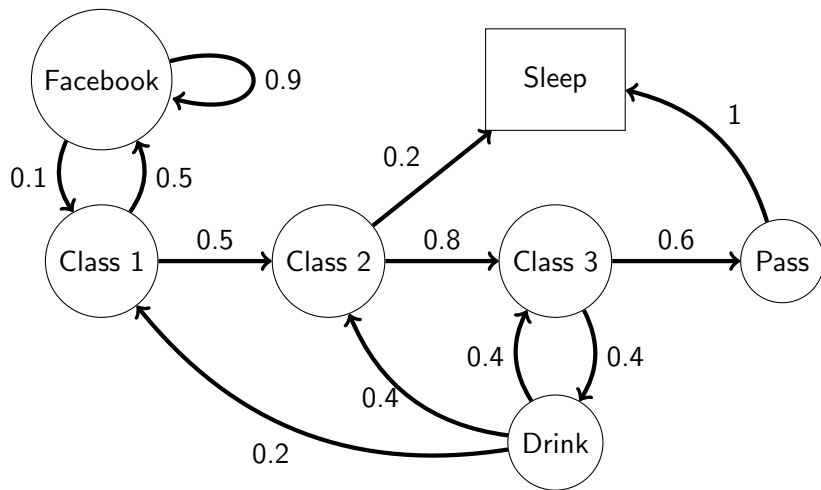
A **Markov process** is a memoryless random process, *i.e.*, a sequence of random states  $S_1, S_2, \dots$  with the Markov property.

## Markov chain

A *Markov chain* is a pair  $(\mathcal{S}, P)$  with

- 1  $\mathcal{S}$  is a finite set of states;
- 2  $P$  is a the transition matrix.

# Student Markov chain




# Student Markov chain episodes

## Sample episodes

$$S_1, S_2, \dots, S_T,$$


starting with  $S_1 = C1$ :

- C1 C2 C3 P S;
- C1 F F F C1 C2 S;
- C1 C2 C3 D C1 C2 C3 D C2 S;
- C1 C2 C3 D C1 F F F F C1 C2 C3 D C2 C3 D C3 P S.






# Student Markov chain transition matrix



	C1	C2	C3	P	D	F	S
C1		0.5				0.5	
C2			0.8				0.2
C3				0.6	0.4		
P							1
D	0.2	0.4	0.4				
F	0.1					0.9	
S							1

# Agent and environment

- The **agent** is the decision maker.
- The **environment** is everything outside the agent.
- These interact continually
  - the agent takes actions;
  - the environment presents new situations and gives rewards.



# Interactions between agent and environment

- At each time step
  - the agent observes the environment's *state*  $S_t \in \mathcal{S}$ ;
  - the agent selects an *action*  $A_t \in \mathcal{A}(S_t)$ ;
  - the agent receives the *reward*  $R_{t+1} \in \mathcal{R}$ ;
  - the agent finds itself in the new state  $S_{t+1} \in \mathcal{S}$ .
- Therefore a trajectory of an MDP is

$$S_0, A_0, R_1, S_1, A_1, R_2, S_2, \dots$$

- The *dynamics* of an MDP are defined by

$$p(s', r | s, a) = \mathbb{P}[S_{t+1} = s', R_{t+1} = r | S_t = s, A_t = a],$$

with

$$\sum_{s' \in \mathcal{S}} \sum_{r \in \mathcal{R}} p(s', r | s, a) = 1, \quad \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}.$$

# Markov property in MDPs

- The state  $s$  of an MDP satisfies the Markov property.
- $\mathbb{P}[S_{t+1}, R_{t+1}]$  depends only on  $S_t$  and  $A_t$ .
- This is an assumption about the representation
  - not the process.
- Markov state can be learned from non-Markov observations.

# Some probability functions

- From  $p(s', r|s, a)$  we can define other probability functions:
  - state-transition probabilities

$$\begin{aligned} p(s'|s, a) &= \mathbb{P}[S_{t+1} = s' | S_t = s, A_t = a] \\ &= \sum_{r \in \mathcal{R}} p(s', r|s, a); \end{aligned}$$

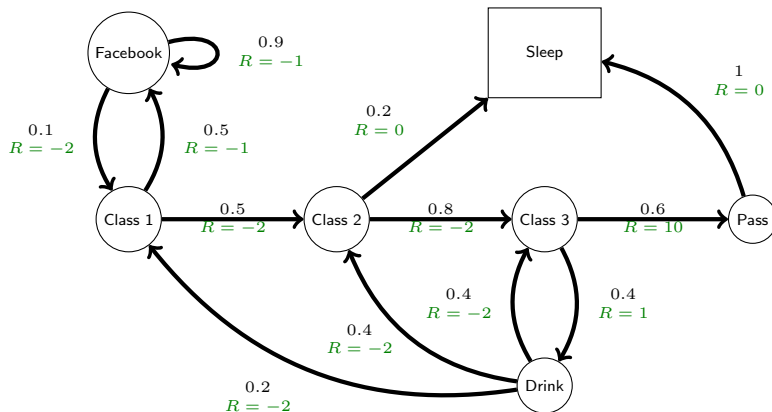
- expected rewards for state-action pairs

$$\begin{aligned} r(s, a) &= \mathbb{E}[R_{t+1} | S_t = s, A_t = a] \\ &= \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r|s, a); \end{aligned}$$

- expected rewards for state-action-next action triples

$$\begin{aligned} r(s, a, s') &= \mathbb{E}[R_{t+1} | S_t = s, A_t = a, S_{t+1} = s'] \\ &= \sum_{r \in \mathcal{R}} r \frac{p(s', r|s, a)}{p(s'|s, a)}. \end{aligned}$$

# Student Markov chain with rewards



# The reward hypothesis

## Reward hypothesis

That all of what we mean by goals and purposes can be well thought of as the maximization of the expected value of the cumulative sum of a received scalar signal (called reward).

Some examples of rewards:

**escape from a maze:**  $-1$  at each time step in the maze;

**completing task:**  $1$  at each step at which the task is completed  
and  $0$  otherwise

**checkers:**  $+1$  for winning and  $-1$  for losing a game.

# Returns

## Episodic tasks

- If there is a natural notion of final time step
  - the agent–environment interaction breaks naturally into subsequences, which we call **episodes**;
  - the time of termination  $T$  is a random variable that varies from episode to episode;
  - each episode ends in a special state called the **terminal state**, followed by a reset;
  - we use  $\mathcal{S}^+$  to denote  $\mathcal{S}$  and the terminal states;
  - the **expected return** is the sum of rewards

$$G_t = R_{t+1} + R_{t+2} + \cdots + R_T.$$



# Returns

## Continuing tasks

- If we are dealing with continuing task
  - introduce a **discount factor**  $\gamma \in [0, 1]$ ;
  - the **expected return** is the sum of discounted rewards

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}. \end{aligned}$$

- delayed rewards are discounted


$$G_t = R_{t+1} + \gamma G_{t+1};$$

- this values immediate reward above delayed reward
  - ▶  $\gamma \rightarrow 0$  leads to greedy evaluation;
  - ▶  $\gamma \rightarrow 1$  leads to far-sighted evaluation;
- discounting with  $\gamma < 1$  avoids infinite returns if  $\mathcal{R}$  is bounded.

# Returns

## Unifying notation

- The terminal state of an episodic task can be thought as an absorbing state generating reward 0.



- With such a convention, the return can be defined as

$$G_t = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

even for episodic tasks.

# Policy


- A policy is a mapping from the current state  $s \in \mathcal{S}$  to probabilities of selecting actions  $a \in \mathcal{A}(s)$ .
- If the agent is following policy  $\pi$

$$\pi(a|s) = \mathbb{P}[A_t = a | S_t = s].$$

- Reinforcement learning methods specify how the agent's policy is changed as a result of its experience.

# Student Markov chain policy

Policy:



$$\begin{aligned}
 \pi(C2|C1) &= 0.5, & \pi(F|C1) &= 0.5, \\
 \pi(C3|C2) &= 0.8, & \pi(S|C2) &= 0.2, \\
 \pi(P|C3) &= 0.6, & \pi(D|C3) &= 0.4, \\
 \pi(S|P) &= 1, \\
 \pi(C1|D) &= 0.2, & \pi(C2|D) &= 0.4, & \pi(C3|D) &= 0.4, \\
 \pi(C1|F) &= 0.1, & \pi(F|F) &= 0.9, \\
 \pi(S|S) &= 1.
 \end{aligned}$$

# Value function

Value functions estimate how good it is for the agent to be in a given state (or state–action pairs).

## State value function

The value function of a state  $s$  under a policy  $\pi$  is the expected return when starting in  $s$  and following  $\pi$  thereafter:

$$v_{\pi}(s) = \mathbb{E}_{\pi}[G_t | S_t = s].$$

## State-action value function

The value function of a state  $s$  and of action  $a$  under a policy  $\pi$  is the expected return starting from  $s$ , taking the action  $a$ , and following policy  $\pi$  thereafter:

$$q_{\pi}(s, a) = \mathbb{E}_{\pi}[G_t | S_t = s, A_t = a].$$

# Bellman equation


- We can obtain a consistency condition for  $v_\pi$ :

$$\begin{aligned}v_\pi(s) &= \mathbb{E}_\pi[G_t | S_t = s] \\&= \mathbb{E}_\pi \left[ \sum_{k=0}^{\infty} \gamma^k R_{t+k+1} | S_t = s \right] \\&= \mathbb{E}_\pi [R_{t+1} + \gamma G_{t+1} | S_t = s] \\&= \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{s' \in \mathcal{S}^+, r \in \mathcal{R}} p(s', r | s, a) (r + \gamma \mathbb{E}_\pi[G_{t+1} | S_{t+1} = s']) \\&= \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{s' \in \mathcal{S}^+, r \in \mathcal{R}} p(s', r | s, a) (r + \gamma v_\pi(s')).\end{aligned}$$

- For each triple  $r, a, s'$ 
  - compute its probability  $\pi(a|s)p(s', r|s, a)$ ;
  - compute the expected return  $r + \gamma v_\pi(s')$ .
- Sum over all possibilities to get an expected value.

# Backup diagram of the Bellman equation

$$v_{\pi}(s) = \sum_{a \in \mathcal{A}(s)} \pi(a|s) \sum_{s' \in \mathcal{S}^+, r \in \mathcal{R}} p(s', r|s, a) (r + \gamma v_{\pi}(s')) .$$



# Relation between $v_\pi$ and $q_\pi$

- It can be easily derived that

- $v_\pi(s) = \sum_{a \in \mathcal{A}(s)} \pi(a|s) q_\pi(s, a);$

- $q_\pi(s, a) = \sum_{s' \in \mathcal{S}, r \in \mathcal{R}} p(s', r|s, a) (r + \gamma v_\pi(s')).$


- This allows us to derive a Bellman equation for  $q_\pi$ :

$$q_\pi(s, a) = \sum_{s' \in \mathcal{S}, r \in \mathcal{R}} p(s', r|s, a) \left( r + \gamma \sum_{a' \in \mathcal{A}(s')} \pi(a'|s') q_\pi(s', a') \right).$$



# Backup diagram of the Bellman equation for $q_\pi$

$$q_\pi(s, a) = \sum_{s' \in \mathcal{S}, r \in \mathcal{R}} p(s', r | s, a) \left( r + \gamma \sum_{a' \in \mathcal{A}(s')} \pi(a' | s') q_\pi(s', a') \right).$$



# Value function for the student Markov chain

## Bellman equation

The Bellman equation for the student Markov chain reads as

$$v_{\pi}(\text{C1}) = 0.5(\gamma v_{\pi}(\text{C2}) - 2) + 0.5(\gamma v_{\pi}(\text{F}) - 1),$$

$$v_{\pi}(\text{C2}) = 0.8(\gamma v_{\pi}(\text{C3}) - 2) + 0.2\gamma v_{\pi}(\text{S}),$$

$$v_{\pi}(\text{C3}) = 0.4(\gamma v_{\pi}(\text{D}) + 1) + 0.6(\gamma v_{\pi}(\text{P}) + 10),$$

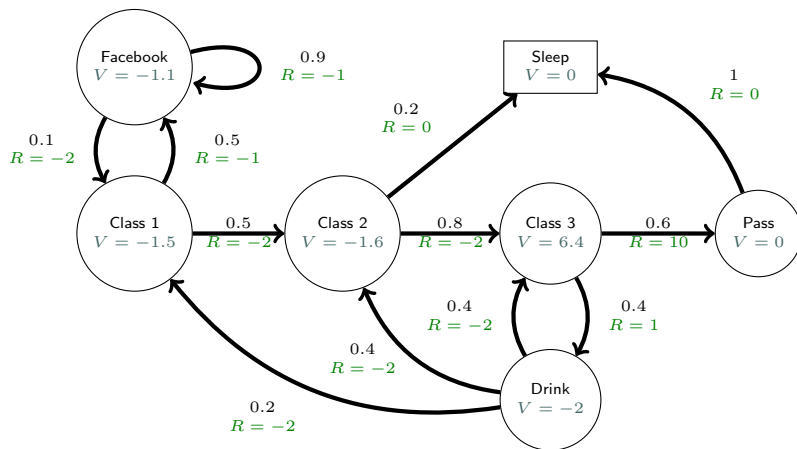
$$v_{\pi}(\text{P}) = \gamma v_{\pi}(\text{S}),$$

$$v_{\pi}(\text{D}) = 0.2(\gamma v_{\pi}(\text{C1}) - 2) + 0.4(\gamma v_{\pi}(\text{C2}) - 2), \\ + 0.4(\gamma v_{\pi}(\text{C3}) - 2),$$

$$v_{\pi}(\text{F}) = 0.1(\gamma v_{\pi}(\text{C1}) - 2) + 0.9(\gamma v_{\pi}(\text{F}) - 1),$$

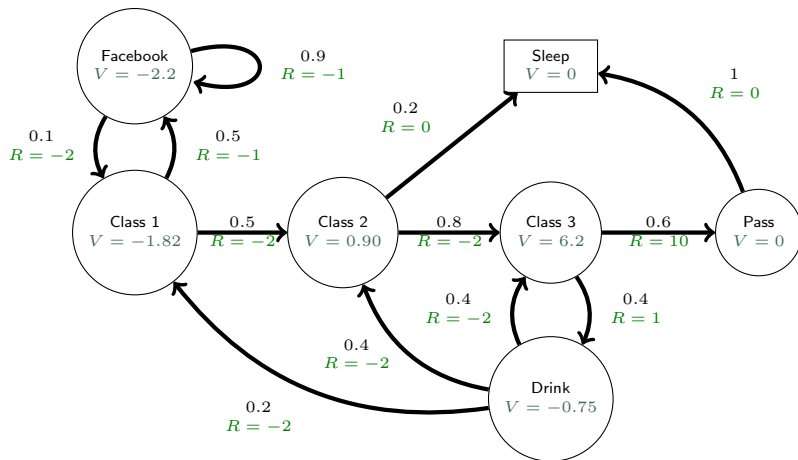
$$v_{\pi}(\text{S}) = \gamma v_{\pi}(\text{S}).$$

# Value function for the student Markov chain

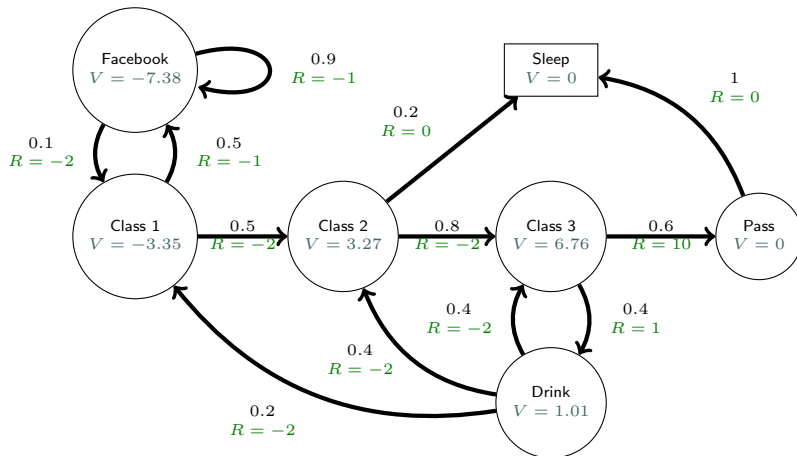
 $\gamma = 0$ 

# Value function for the student Markov chain

$\gamma = 0.5$



# Value function for the student Markov chain

 $\gamma = 0.9$ 

# Optimal value function

- Value functions can be used to sort policies
  - $\pi \geq \pi' \iff v_\pi(s) \geq v_{\pi'}(s), \forall s \in \mathcal{S}.$
- There is always at least one policy that is better than or equal to all other policies
  - this is an optimal policy, denoted  $\pi_*$ ;
  - all policies  $\pi_*$  share the same value function
    - ▶ this is the **optimal value function**

$$v_*(s) = \max_{\pi} v_\pi(s);$$

- ▶ all policies  $\pi_*$  share the same optimal action-value function

$$\begin{aligned} q_*(s, a) &= \max_{\pi} q_\pi(s, a) \\ &= \mathbb{E}[R_{t+1} + \gamma v_*(S_{t+1}) | S_t = s, A_t = a]. \end{aligned}$$

# Bellman optimality equation

- $v_*$  and  $q_*$  must satisfy the Bellman equation.
- Further, it must hold that



$$v_*(s) = \max_{a \in \mathcal{A}(s)} q_*(s, a).$$

- We thus obtain the **Bellman optimality equation**

$$v_*(s) = \max_{a \in \mathcal{A}(s)} \sum_{s' \in \mathcal{S}^+, r \in \mathcal{R}} p(s', r | s, a) (r + \gamma v_*(s')),$$

$$q_*(s, a) = \sum_{s' \in \mathcal{S}, r \in \mathcal{R}} p(s', r | s, a) \left( r + \gamma \max_{a' \in \mathcal{A}(s')} q_*(s', a') \right).$$

# Backup diagrams of the Bellman optimality equation





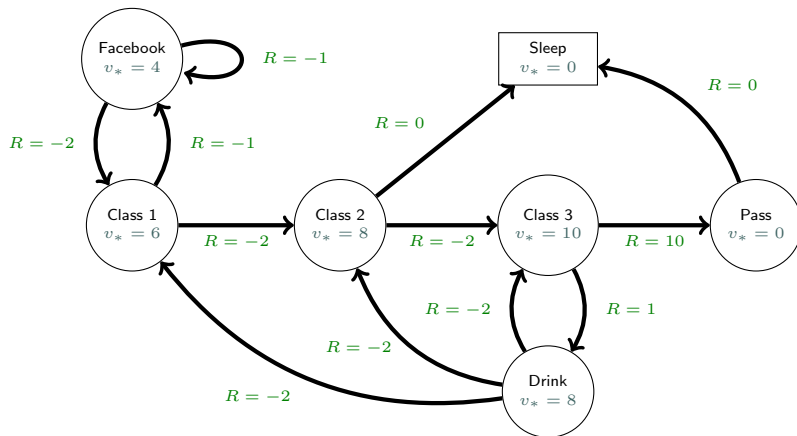
# Considerations on the Bellman optimality equation

- The Bellman optimality equation
  - is nonlinear;
  - no closed form solution exists;
  - can be solved explicitly in some cases
    - ▶ Dijkstra's algorithm;
    - ▶ A\* search algorithm;
- Optimal actions at state  $s$  can be determined as

$$a \leftarrow \arg \max_a q_*(s, a)$$

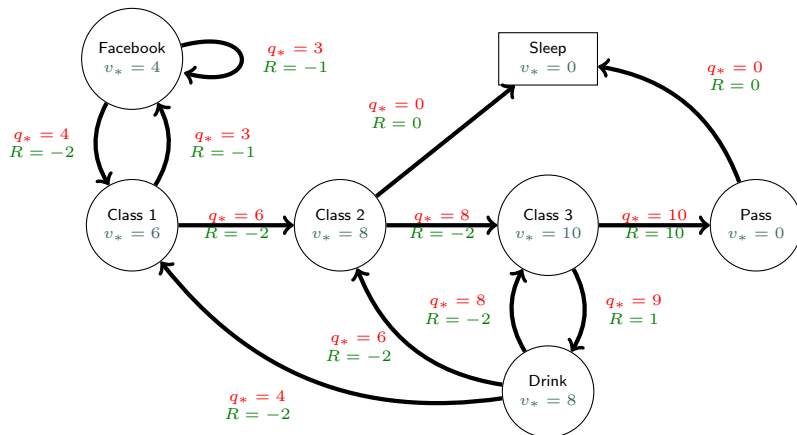
- there is always a deterministic optimal policy for any MDP.

# Optimal value function in student MDP

 $\gamma=1$ 

# Optimal state-action value function in student MDP

$\gamma=1$



# Issues on solving the Bellman optimality equation

- The dynamics of the environment are not accurately known;
- Computationally expensive;
- The states may not have the Markov property.
- Many iterative solution methods
  - value iteration;
  - policy iteration;
  - q-learning;
  - SARSA.