

# Temporal-difference learning

Corrado Possieri

Machine and Reinforcement Learning in Control Applications

# Introduction

- Learn directly from experience without a model of the environment.
- As DP, use learned estimates to update the prediction
  - learns from incomplete episodes;
  - bootstrap.
- Updates a guess towards a guess.
- Can be used both for prediction and control.

# Monte Carlo vs temporal-difference prediction

- Given a policy  $\pi$ , the goal is to estimate  $v_\pi$ .
- MC methods wait until the return following the visit is known
  - use the return  $G_t$  as a target for  $V(S_t)$

$$V(S_t) \leftarrow V(S_t) + \alpha(G_t - V(S_t)).$$

- The simplest TD method (TD(0)) uses the immediate reward
  - use the reward  $R_{t+1}$
  - and the expected return  $\gamma V(S_{t+1})$  as a target for  $V(S_t)$

$$V(S_t) \leftarrow V(S_t) + \alpha(R_{t+1} + \gamma V(S_{t+1}) - V(S_t)).$$

# One-step temporal-difference

## Tabular TD(0)

**Input:**  $\pi$

**Output:**  $v_\pi$

### Initialization

$V(s) \leftarrow \text{arbitrary}, \forall s \in \mathcal{S}$

$V(\text{terminal}) \leftarrow 0$

### Loop

initialize  $S$

**for** each step of the episode **do**

$A \leftarrow \pi(S)$

take action  $A$  and observe  $R, S'$

$V(S) \leftarrow V(S) + \alpha(R + \gamma V(S') - V(S))$

$S \leftarrow S'$

**if**  $S$  is terminal **then**

reinitialize the episode

# Comparison of DP, MC, and TD

- DP estimates  $v_\pi(s)$  by bootstrapping

$$V(s) \leftarrow \mathbb{E}_\pi[R_{t+1} + \underbrace{\gamma V(S_{t+1})}_{\text{estimate}} | S_t = s].$$

- MC estimates  $v_\pi(s)$  by sample mean

$$V(s) \leftarrow \underbrace{\mathbb{E}_\pi}_{\text{sample}} [G_t | S_t = s].$$

- TD estimates  $v_\pi(s)$  by estimating both

$$V(s) \leftarrow \underbrace{\mathbb{E}_\pi}_{\text{sample}} [R_{t+1} + \underbrace{\gamma V(S_{t+1})}_{\text{estimate}} | S_t = s].$$

# Temporal-difference error


- The *TD error* is

$$\delta_t = R_{t+1} + \gamma V(S_{t+1}) - V(S_t).$$

- This error is used to update  $V(S_t)$  towards a better estimate.
- It is not available until  $t + 1$ .
- If  $V$  does not change, the MC error satisfies

$$\begin{aligned} G_t - V(S_t) &= R_{t+1} + \gamma G_{t+1} - V(S_t) \\ &= \delta_t + \gamma(G_{t+1} - V(S_{t+1})) \\ &= \delta_t + \gamma\delta_{t+1} + \gamma^2(G_{t+1} - V(S_{t+1})) \\ &= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k. \end{aligned}$$

# Unified view





# Driving home example

State	Elapsed time	Predicted time to go	Predicted total time
Leaving university	0	30	30
Reaching car	5	35	40
Exiting highway	20	15	35
Secondary road	30	10	40
Entering home	40	3	43
Home	43	0	43





# MC vs TD updates

## MC updates



## TD updates



# Advantages of TD

- TD can be implemented online.
- MC waits until the end of the episode.
- TD require just a single step.
- MC requires complete sequences.
- TD works also for incomplete sequences.
- MC can be applied just for episodic tasks.
- TD can be used for continuous and episodic tasks.

# Bias and variance of estimators

- MC target

$$G_t = R_{t+1} + \gamma R_{t+2} + \dots + \gamma^{T-1} R_T$$

is an unbiased estimate of  $v_\pi(S_t)$ .

- TD true target

$$R_{t+1} + \gamma v_\pi(S_{t+1})$$

is an unbiased estimate of  $v_\pi(S_t)$ .

- TD target

$$R_{t+1} + \gamma V(S_{t+1})$$

is a biased estimate of  $v_\pi(S_t)$ .

- TD target has lower variance than MC target
  - return depends on *many* random actions, transitions, rewards;
  - TD target depends on *one* random action, transition, reward.

# Advantages and disadvantages of MC and TD

- MC has high variance and zero bias
  - good convergence properties;
  - not very sensitive to initial value;
  - simple to understand and use;
  - more efficient in non-Markov environments.
- TD has low variance and nonzero bias
  - usually more efficient than MC;
  - TD(0) proved to converge to  $v_\pi$ 
    - ▶ in the mean for constant (small)  $\alpha$ ;
    - ▶ almost surely if  $\sum_k \alpha_k = \infty$  and  $\sum_k \alpha_k^2 < \infty$ ;
  - sensitive to initial value;
  - more efficient in Markov environments.

# Batch updating


- Suppose to have a finite amount of experience.
- We can present the experience repeatedly until convergence.
- Given  $V$ , update it only once for each batch
  - compute the updates  $\alpha(\text{target}_t - V(S_t))$  at each time step;
  - change the value function once with the sum of all increments.
- Constant  $\alpha$  TD converges deterministically.
- Constant  $\alpha$  MC converges deterministically.
- The two results may be different.

# You are the predictor (1)

- Suppose you observe the following eight episodes
  - 1  $A, 0, B, 0;$
  - 2  $B, 1;$
  - 3  $B, 1;$
  - 4  $B, 1;$
  - 5  $B, 1;$
  - 6  $B, 1;$
  - 7  $B, 1;$
  - 8  $B, 0.$
- We want to estimate  $V(A)$  and  $V(B)$ .

## You are the predictor (2)

- The optimal value for  $B$  is  $V(B) = \frac{6}{8} = 0.75$ .
- Modeling experience via an MP



- $V(A) = 0.75$ ;
  - same answer given by TD(0).
- We observed the return from  $A$  once and it was 0
    - $V(A) = 0$ ;
    - same answer given by MC;
    - minimum squared error on training data.

# Certainty equivalence

- Batch MC converges to solution with minimum MS error
  - best fit to the observed returns

$$\sum_{k=1}^K \sum_{t=1}^{T_k} (G_t^k - V(S_t^k))^2.$$

- Batch TD(0) converges to solution of max likelihood MDP
  - solution to the MDP that best fits the data

$$\hat{P}_{s,s',a} = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_{t+1}^k = s', s_t^k = s, a_t^k = a),$$

$$\hat{R}_{s,a} = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} r_t^k \mathbf{1}(s_t^k = s, a_t^k = a);$$

- equivalent to assuming that the process estimate was known.



# Off-policy TD prediction

- Use TD targets generated from  $b$  to evaluate  $\pi$ .
- Weight TD target by importance sampling.
- Only need a single importance sampling correction

$$V(S_t) \leftarrow V(S_t) + \alpha \left( \frac{\pi(A_t|S_t)}{b(A_t|S_t)} \underbrace{(R_{t+1} + \gamma V(S_{t+1}))}_{\text{TD target}} - V(S_t) \right).$$

- Lower variance than Monte-Carlo importance sampling.

# On-policy TD control

- We follow the path of GPI.
- Use TD for the evaluation part.
- Estimate  $q_\pi$  rather than  $v_\pi$ .
- Transitions from a state-action pair to a state-action pair.



- This is still an MDP.

# SARSA

- Apply TD(0) to action values

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha \underbrace{\underbrace{(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}))}_{\text{TD target}} - Q(S_t, A_t)}_{\text{TD error } \delta_t}.$$

- If  $S_{t+1}$  is terminal

$$Q(S_{t+1}, A_{t+1}) = 0, \quad \forall A_t.$$

- The update is carried out on the basis of the quintuple

$$(S_t, A_t, R_{t+1}, S_{t+1}, A_{t+1}).$$

# Greedy in the Limit with Infinite Exploration

## GLIE policies

A policy  $\pi$  is GLIE if

- 1 All state-action pairs are explored infinitely many times

$$N_k(s, a) \rightarrow \infty, \quad \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}.$$

- 2 The policy converges on a greedy policy

$$\pi(a|s) = \mathbf{1}(a = \arg \max_a Q(s, a)).$$

- For instance,  $\varepsilon$ -greedy policies are GLIE if  $\varepsilon_k = \frac{1}{k}$ .

# Notes on SARSA


- If  $Q$  does not change, the MC error satisfies

$$\begin{aligned} G_t - Q(S_t, A_t) &= R_{t+1} + \gamma G_{t+1} - Q(S_t, A_t) \\ &= \delta_t + \gamma(G_{t+1} - Q(S_{t+1}, A_{t+1})) \\ &= \delta_t + \gamma\delta_{t+1} + \gamma^2(G_{t+2} - Q(S_{t+2}, A_{t+2})) \\ &= \sum_{k=t}^{T-1} \gamma^{k-t} \delta_k. \end{aligned}$$

- If the policy is GLIE and

$$\sum_k \alpha_k = \infty, \quad \sum_k \alpha_k^2 < \infty$$

- SARSA converges with probability 1.



# On-policy TD control algorithm

## SARSA algorithm

**Input:**  $\alpha > 0, \varepsilon > 0$

**Output:**  $q_*, \pi_*$

### Initialization

$Q(s, a) \leftarrow \text{arbitrary}, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}$

$Q(\text{terminal}, \cdot) \leftarrow 0$

### Loop

initialize  $S$

$A \leftarrow$  action derived by  $Q(S, \cdot)$  (e.g.,  $\varepsilon$ -greedy)

**for** each step of the episode **do**

take action  $A$  and observe  $R, S'$

$A' \leftarrow$  action derived by  $Q(S', \cdot)$  (e.g.,  $\varepsilon$ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma Q(S', A') - Q(S, A))$

$S \leftarrow S'$

$A \leftarrow A'$

**if**  $S$  is terminal **then**

reinitialize the episode

# Q-learning

- Independent of the policy being followed.
- Directly approximate  $q_*$ .
- Update  $Q$  as

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \max_a Q(S_{t+1}, a) - Q(S_t, A_t)).$$


- No importance sampling is required.

# Notes on Q-learning

- Next action is chosen using behavior policy, e.g.,  $\varepsilon$ -greedy.
- We consider alternative successor action following the greedy target policy  $\pi$ .
- Both behavior and target policies improve.
- If all pairs continue to be updated and

$$\sum_k \alpha_k(x_k, a_k) = \infty, \quad \sum_k \alpha_k^2(x_k, a_k) < \infty$$

- Q-learning converges with probability 1.





# Off-policy TD control algorithm

## Q-learning algorithm

**Input:**  $\alpha > 0, \varepsilon > 0$

**Output:**  $q_*, \pi_*$

### Initialization

$Q(s, a) \leftarrow \text{arbitrary}, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}$

$Q(\text{terminal}, \cdot) \leftarrow 0$

### Loop

initialize  $S$

**for** each step of the episode **do**

$A \leftarrow$  action derived by  $Q(S, \cdot)$  (e.g.,  $\varepsilon$ -greedy)

take action  $A$  and observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_a Q(S', a) - Q(S, A))$


$S \leftarrow S'$

**if**  $S$  is terminal **then**

reinitialize the episode


# Cliff walking

- Consider the gridworld on the right.
- Terminal states are those in red and green.
- Reward is  $-1$  on all transitions except those into the red area.
- Stepping into this region incurs a reward of  $-100$ .




# Comparison of SARSA and Q-learning on cliff walking

SARSA - episode 1000



Q-learning - episode 1000



# Expected SARSA


- Consider the basic update of Q-learning.
- Rather than maximizing, take expectation

$$\begin{aligned} Q(S_t, A_t) &\leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma \mathbb{E}_{\pi}[Q(S_{t+1}, A_{t+1})|S_t] - Q(S_t, A_t)) \\ &= Q(S_t, A_t) + \alpha \left( R_{t+1} + \gamma \left( \sum_a \pi(a|S_{t+1}) Q(S_{t+1}, a) \right) - Q(S_t, A_t) \right). \end{aligned}$$

- This moves deterministically as SARSA moves in expectation.
- It eliminates the variance due to the random selection of  $A_{t+1}$ .

# Notes on expected SARSA

- Retains advantages of SARSA.
- Consistent empirical advantage of expected SARSA over SARSA.
- Can select large values of  $\alpha$  ( $\simeq 1$ ) when the environment is non-stochastic.
- Can be used off-policy
  - $\pi$  is the greedy policy
  - $b$  is the behavior policy
    - expected SARSA is Q-learning.



# Maximization bias

- All control algorithms discussed so far involve maximization
  - in Q-learning the target policy is the greedy policy on  $Q$ ;
  - in SARSA in Sarsa the policy is often  $\varepsilon$ -greedy.
- A maximization over estimates may lead to a (positive) bias.
- Consider, e.g., the case
  - $q(s, a) = 0$  for all  $a \in \mathcal{A}(s)$ ;
  - $Q(s, a)$  are uncertain and thus distributed some above and some below zero.
- The problem is due to the fact that the same data are used to estimate both optimal actions and their values.

# Double learning

- Divide data in two sets and used them to learn two independent estimates:  $Q_1$  and  $Q_2$ .

- $Q_1$  can be used to determine the maximizing action

$$A_t^* = \arg \max_a Q_1(S_t, a).$$

- $Q_2$  can be used to estimate its value

$$Q_2(S_t, A_t^*) = Q_2(S_t, \arg \max_a Q_1(S_t, a)).$$

- This is equivalent to

$$Q_2(s, a) \leftarrow Q_2(s, a) + \alpha(R + \gamma Q_1(s', \arg \max_{a'} Q_2(s', a')) - Q_2(s, a)).$$

- The estimate will be unbiased

$$\mathbb{E}[Q_2(S_t, A_t^*)] = q(S_t, A_t^*).$$

# Double Q-learning

## Double Q-learning algorithm

**Input:**  $\alpha > 0, \varepsilon > 0$

**Output:**  $q_*, \pi_*$

### Initialization

$Q_1(s, a) \leftarrow \text{arbitrary}, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}, Q_1(\text{terminal}, \cdot) \leftarrow 0$

$Q_2(s, a) \leftarrow \text{arbitrary}, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}, Q_2(\text{terminal}, \cdot) \leftarrow 0$

### Loop

initialize  $S$

**for** each step of the episode **do**

$A \leftarrow$  action derived by  $Q_1(S, \cdot) + Q_2(S, \cdot)$  (e.g.,  $\varepsilon$ -greedy)

take action  $A$  and observe  $R, S'$

with probability 0.5

$Q_1(S, A) \leftarrow Q_1(S, A) + \alpha(R + \gamma Q_2(S', \arg \max_a Q_1(S', a)) - Q_1(S, A))$

else

$Q_2(S, A) \leftarrow Q_2(S, A) + \alpha(R + \gamma Q_1(S', \arg \max_a Q_2(S', a)) - Q_2(S, A))$

$S \leftarrow S'$

**if**  $S$  is terminal **then**

reinitialize the episode