

# Coupling planning and learning

Corrado Possieri

Machine and Reinforcement Learning in Control Applications

# Introduction

Planning: model-based.

Learning: model free.

- The hearth of both methods is computing a value function
  - look ahead to future events;
  - compute a backed-up value;
  - update target.
- Can these methods be intermixed?


# Model, learning, and planning

- A model is anything the agent can use to predict the behavior of the environment

**distribution models:** gives all possibilities and their probability;

**sample models:** produce just one of the possibilities.

- Distribution models can be used to generate samples.



- Learning can be used for planning.


# Sample-based planning

- Use the model only to generate samples.
- Sample experience from model

$$S_{t+1} \sim p(S_{t+1}|S_t, A_t),$$
$$R_{t+1} = r(S_t, A_t).$$

- Apply model-free RL to samples.
- Often more efficient than planning.

# Indirect reinforcement learning



# Notes on indirect learning

- Indirect methods often make fuller use of a limited amount of experience.
- Direct methods are simpler and are not affected by biases.
- Model-based RL is only as good as the estimated model.
- When the model is inaccurate, planning will compute a suboptimal policy.

# How to learn a model

- In deterministic environments
  - in state  $S_t$ , take action  $A_t$ ;
  - observe  $R_{t+1}$  and  $S_{t+1}$ ;
  - $p(S_{t+1}|S_t, A_t) \leftarrow 1$ ,  $p(s|S_t, A_t) \leftarrow 0, \forall s \neq S_{t+1}$ ;
  - $r(S_t, A_t) \leftarrow R_{t+1}$ .
- In MDP environments
  - observing the history
    - ▶ determining  $p(s'|s, a)$  is a *density estimation* problem;
    - ▶ determining  $r(s, a)$  is a *regression* problem.

# Table lookup model

- Assuming  $\mathcal{S}$  and  $\mathcal{A}$  known
  - for and MDP we need to estimate  $P$  and  $R$ 
    - ▶ use empirical samples

$$\hat{P}_{s,s',a} = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} \mathbf{1}(s_{t+1}^k = s', s_t^k = s, a_t^k = a),$$

$$\hat{R}_{s,a} = \frac{1}{N(s,a)} \sum_{k=1}^K \sum_{t=1}^{T_k} r_t^k \mathbf{1}(s_t^k = s, a_t^k = a).$$



# Dyna-Q

- Use deterministic model learning
- Use one-step Q-learning as planning method
  - randomly sample from state-action pairs that have been previously experienced;
  - the model returns the last-observed next state and next reward as its prediction.
- Use one-step Q-learning as direct RL method.
- Learning and planning differ only in the source of their experience.

# Dyna-Q algorithm

## Dyna-Q algorithm

**Input:**  $\alpha > 0, \varepsilon > 0$

**Output:**  $q_*, \pi_*$

### Initialization

$Q(s, a) \leftarrow \text{arbitrary}, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}; Q(\text{terminal}, \cdot) \leftarrow 0$

$\text{Model}(s, a) \leftarrow \emptyset, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}$

### Loop

$S \leftarrow \text{current state}$

$A \leftarrow \varepsilon - \text{greedy}(S, Q)$

take action  $A$  and observe  $R, S'$

$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_a Q(S', a) - Q(S, A))$

$\text{Model}(S, A) \leftarrow S', R$

### repeat

$(S, A) \leftarrow \text{random previously experienced pair}$

$R, S' \leftarrow \text{Model}(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_a Q(S', a) - Q(S, A))$

**until** next sample is available

# Notes on Dyna-Q

- Learns much faster in deterministic environments.
- If the environment changes, it can adapt.
- However, the formerly correct policy may not reveal improvements.
- The planning process is likely to compute a suboptimal policy.
- Exploration/exploitation conflict in a planning context
  - exploration: try to improve the model;
  - exploitation: take the best possible action according to the available model.

# Dyna-Q+

- Keep track for each state–action pair of how many time steps have elapsed since last visit.
- The more time that has elapsed, the more is like that the model is incorrect.
- Encourage behavior that tests long-untried actions
  - the modeled reward for a transition is  $r$ ;
  - the transition has not been tried in  $\tau$  time steps;
  - planning assumes that the reward is  $r + \kappa\sqrt{\tau}$ ;
  - an alternative is to select action as that maximizing

$$Q(S_t, a) + \kappa\sqrt{\tau(S_t, a)}.$$

# Prioritized sweeping

- Planning is more efficient if simulated transitions and updates are focused on particular state–action pairs.
- If simulated transitions are generated uniformly, then many wasteful updates are made.
- In general, we want to work back from any state whose value has changed
  - the predecessor pairs of those that have changed are more likely to also change;
  - prioritize the updates according to a measure of their urgency;
  - needs an inverse model.

# Prioritized sweeping algorithm

## Prioritized sweeping algorithm

**Input:**  $\alpha > 0$ ,  $\varepsilon > 0$ , threshold  $\theta > 0$

**Output:**  $q_*$ ,  $\pi_*$

### Initialization

$Q(s, a) \leftarrow \text{arbitrary}, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}; Q(\text{terminal}, \cdot) \leftarrow 0$

$\text{Model}(s, a) \leftarrow \emptyset, \forall a \in \mathcal{A}(s), \forall s \in \mathcal{S}$

$\text{PQueue} \leftarrow \emptyset$

### Loop

$S \leftarrow \text{current state}$

$A \leftarrow \varepsilon\text{-greedy}(S, Q)$

take action  $A$  and observe  $R, S'$

$\text{Model}(S, A) \leftarrow S', R$

$P \leftarrow |R + \gamma \max_a Q(S', a) - Q(S, A)|$

**if**  $P > \theta$  **then**

    insert  $S, A$  into PQueue with priority  $P$

**while**  $\text{PQueue} \neq \emptyset$  **do**

$(S, A) \leftarrow \text{first}(\text{PQueue})$

$R, S' \leftarrow \text{Model}(S, A)$

$Q(S, A) \leftarrow Q(S, A) + \alpha(R + \gamma \max_a Q(S', a) - Q(S, A))$

**for all**  $\bar{S}, \bar{A}$  predicted to lead to  $S$  **do**

$\bar{R}, \bar{S}' \leftarrow \text{Model}(\bar{S}, \bar{A})$

$\bar{P} \leftarrow |\bar{R} + \gamma \max_a Q(\bar{S}', a) - Q(\bar{S}, \bar{A})|$

**if**  $\bar{P} > \theta$  **then**

            insert  $\bar{S}, \bar{A}$  into PQueue with priority  $\bar{P}$

# Trajectory sampling

- Distribute updates according to the on-policy distribution
  - distribution observed when following the current policy;
  - simulate individual trajectories and perform updates at the state encountered along the way.
- States actually visited are updated more often.
- Uninteresting parts of the space are ignored.
- This is the same that happens in real time DP
  - can find a policy that is optimal on the relevant states without visiting every state infinitely often.


# Planning at decision time

- Planning executed at current state.
- The values and policy are specific to the current state.
- The values and policy created are discarded after being used.
- Useful in applications in which fast responses are not required



# Heuristic search

- Go through the tree of possible continuations.
- Use a model of the sub-MDP starting from *now*.
- Focused on state/actions that immediately follow.
- Build a search tree with  $S_t$  at its root.



# Rollout algorithms

- Heuristic search guided by MC simulation.
- Simulate episodes from now with the model.
- Average returns of simulated trajectories that start with each action and then follow rollout policy
  - simulate  $K$  episodes following first action  $a$  and then policy  $\pi$

$$\textcolor{red}{S}_t, \textcolor{red}{a}, R_{t+1}^k, S_{t+1}^k, A_{t+1}^k, R_{t+2}^k, \dots, R_T^k, S_T^k;$$

- evaluate actions by mean return

$$Q(\textcolor{red}{S}_t, \textcolor{red}{a}) = \frac{1}{K} \sum_{k=1}^K G_t^k;$$

- take action that maximize  $Q$

$$A_t = \arg \max_a Q(S_t, a).$$

# Monte Carlo tree search

- Use rollout method by modifying policy.
- Record the values of  $Q$  in the search tree.
- In the tree, we pick actions to maximize  $Q$  (e.g.,  $\epsilon$ -greedy).
- Outside the tree use a default policy.
- MC control applied to simulated experience.
- Expand the part of the tree that looks promising.
- MC can be substituted by TD.

# Monte Carlo tree search logic

