

Control with function approximation

Corrado Possieri

Machine and Reinforcement Learning in Control Applications

Introduction

- We now want to build and approximate

$$\hat{q}(s, a, \mathbf{w}) \simeq q_*(s, a).$$

- Natural extension of prediction in the episodic case.
- Attention needed in the continuing case.
- We follow the general pattern of on-policy GPI.

Episodic Semi-gradient Control

- The target update U_t can be any approximation of q_π
 - MC update

$$S_t, A_t \mapsto G_t;$$

- TD(0) update (SARSA)

$$S_t, A_t \mapsto R_{t+1} + \gamma Q(S_{t+1}, A_{t+1});$$

- n -step TD update

$$S_t, A_t \mapsto G_{t:t+n}.$$

- The general update form is

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha(U_t - \hat{q}(S_t, A_t, \mathbf{w}_t))\nabla\hat{q}(S_t, A_t, \mathbf{w}_t).$$

- Couple action-value prediction methods with techniques for policy improvement and action selection.

Semi-gradient SARSA algorithm

Semi-gradient SARSA algorithm

Input: $\alpha > 0$, $\varepsilon > 0$, approximation function \hat{q}

Output: approximate of q_* and π_*

Initialization

$\mathbf{w} \leftarrow$ arbitrarily

Loop

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

for each step of the episode **do**

take action A and observe R, S'

if S' is terminal **then**

$\mathbf{w} \leftarrow \mathbf{w} + \alpha(R - \hat{q}(S, A, \mathbf{w}))\nabla\hat{q}(S, A, \mathbf{w})$

reinitialize the episode

choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$\mathbf{w} \leftarrow \mathbf{w} + \alpha(R + \gamma\hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w}))\nabla\hat{q}(S, A, \mathbf{w})$

$S \leftarrow S'$

$A \leftarrow A'$

n -step semi-gradient SARSA algorithm

n -step semi-gradient SARSA algorithm

Input: $\alpha > 0$, a positive integer n , approximation function \hat{q}

Output: approximate of q_* and π_*

Initialization

$\mathbf{w} \leftarrow$ arbitrary

Loop

initialize $S_0 \neq$ terminal

store $A_0 \leftarrow \varepsilon$ -greedy($\hat{q}(S_0, \cdot, \mathbf{w})$)

$T \leftarrow \infty$

for $t = 0, 1, 2, \dots$ **do**

take action A_t

observe and store R_{t+1} and S_{t+1}

if S_{t+1} is terminal **then**

$T \leftarrow t + 1$

else

store $A_{t+1} \leftarrow \varepsilon$ -greedy($\hat{q}(S_{t+1}, \cdot, \mathbf{w})$)

$\tau = t - n + 1$

if $\tau \geq 0$ **then**

$G \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} \gamma^{i-\tau-1} R_i$

if $\tau + n < T$ **then**

$G \leftarrow G + \gamma^n \hat{q}(S_{\tau+n}, A_{\tau+n}, \mathbf{w})$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha(G - \hat{q}(S_\tau, A_\tau, \mathbf{w})) \nabla \hat{q}(S_\tau, A_\tau, \mathbf{w})$

if $\tau = T - 1$ **then**

proceed to next episode

Average return

- The average reward setting applies to continuing problems.
- No discounting (delayed rewards count as immediate reward).
- The quality of π is defined as the average rate of reward

$$r(\pi) = \lim_{h \rightarrow \infty} \frac{1}{h} \sum_{t=1}^h \mathbb{E}[R_t | S_0, A_{0:t} \sim \pi].$$

- If the MDP is *ergodic*, i.e., the steady state distribution

$$\mu_\pi(s) = \lim_{t \rightarrow \infty} \mathbb{P}[S_t = s | A_{0:t} \sim \pi]$$

exists and is independent of S_0 , then

$$\begin{aligned} r(\pi) &= \lim_{t \rightarrow \infty} \mathbb{E}[R_t | S_0, A_{0:t} \sim \pi] \\ &= \sum_s \mu_\pi(s) \sum_a \pi(a|s) \sum_{s', r} p(s', r | s, a) r. \end{aligned}$$

Differential return and differential value functions

- In the average reward case, we consider the *differential return*

$$G_t = R_{t+1} - r(\pi) + R_{t+2} - r(\pi) + R_{t+3} - r(\pi) + \dots$$

- *Differential value functions* are defined upon G_t

$$v_\pi(s) = \mathbb{E}[G_t | S_t = s],$$
$$q_\pi(s, a) = \mathbb{E}[G_t | S_t = s, A_t = a].$$

- The corresponding Bellman equations are

$$v_\pi(s) = \sum_a \pi(a|s) \sum_{s',r} p(s', r|s, a) (r - r(\pi) + v_\pi(s')),$$
$$q_\pi(s, a) = \sum_{s',r} p(s', r|s, a) (r - r(\pi) + \sum_{a'} \pi(a'|s') q_\pi(s', a')),$$
$$v_*(s) = \max_a \sum_{s',r} p(s', r|s, a) (r - \max_\pi r(\pi) + v_*(s')),$$
$$q_*(s, a) = \sum_{s',r} p(s', r|s, a) (r - \max_\pi r(\pi) + \max_{a'} q_*(s', a')).$$

Differential errors

- The TD differential errors are

$$\delta_t = R_{t+1} - \bar{R}_t + \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t),$$

$$\delta_t = R_{t+1} - \bar{R}_t + \hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t),$$

$$\delta_t = R_{t+1} - \bar{R}_{t+n-1} + R_{t+2} - \bar{R}_{t+n-1} + \cdots + R_{t+n} - \bar{R}_{t+n-1} \\ + \hat{q}(S_{t+n}, A_{t+n}, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t),$$

where \bar{R}_t is an estimate of $r(\pi)$.

- With these definitions, we can implement most algorithms
 - e.g., the SARSA update is

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha \delta_t \nabla \hat{q}(S_t, A_t, \mathbf{w}_t).$$

- ▶ converges to a differential values plus an arbitrary offset;
- ▶ the Bellman equations and the TD errors are unaffected if all the values are shifted by the same amount.

Differential semi-gradient SARSA

Differential semi-gradient SARSA algorithm

Input: $\alpha > 0$, $\beta > 0$, $\varepsilon > 0$, approximation function \hat{q}

Output: approximate of q_* and π_*

Initialization

$\mathbf{w} \leftarrow$ arbitrarily

$\bar{R} \leftarrow$ arbitrarily

Loop

$S, A \leftarrow$ initial state and action of episode (e.g., ε -greedy)

for each step of the episode **do**

take action A and observe R, S'

choose A' as a function of $\hat{q}(S', \cdot, \mathbf{w})$ (e.g., ε -greedy)

$\delta \leftarrow (R - \bar{R} + \hat{q}(S', A', \mathbf{w}) - \hat{q}(S, A, \mathbf{w}))$

$\bar{R} \leftarrow \bar{R} + \beta \delta$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta \nabla \hat{q}(S, A, \mathbf{w})$

$S \leftarrow S'$

$A \leftarrow A'$

Differential n -step semi-gradient SARSA algorithm

Differential n -step semi-gradient SARSA algorithm

Input: $\alpha > 0$, $\beta > 0$, a positive integer n , approximation function \hat{q}

Output: approximate of q_* and π_*

Initialization

$\mathbf{w} \leftarrow$ arbitrary, $\bar{R} \leftarrow$ arbitrary

Loop

initialize $S_0 \neq$ terminal

store $A_0 \leftarrow \varepsilon$ -greedy($\hat{q}(S_0, \cdot, \mathbf{w})$)

$T \leftarrow \infty$

for $t = 0, 1, 2, \dots$ **do**

take action A_t , observe and store R_{t+1} and S_{t+1}

if S_{t+1} is terminal **then**

$T \leftarrow t + 1$

else

store $A_{t+1} \leftarrow \varepsilon$ -greedy($\hat{q}(S_{t+1}, \cdot, \mathbf{w})$)

$\tau = t - n + 1$

if $\tau \geq 0$ **then**

$\delta \leftarrow \sum_{i=\tau+1}^{\min(\tau+n, T)} (R_i - \bar{R}) - \hat{q}(S_\tau, A_\tau, \mathbf{w})$

if $\tau + n < T$ **then**

$\delta \leftarrow \delta + \hat{q}(S_{\tau+n}, A_{\tau+n}, \mathbf{w})$

$\bar{R} \leftarrow \bar{R} + \beta \delta$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta \nabla \hat{q}(S_\tau, A_\tau, \mathbf{w})$

if $\tau = T - 1$ **then**

proceed to next episode

Deprecating discounted setting

- Averaging the rewards over a long interval leads to the average reward setting.
 - the average of discounted returns equals $\frac{r(\pi)}{1-\gamma}$.
- The value of γ has no effect with function approximation.
- We lost the policy improvement theorem using function approximation.
- Discounting algorithms with function approximation do not optimize discounted value over the on-policy distribution,

SARSA(λ)

- Eligibility traces can be used also for control.
- The off-line λ -return algorithm uses \hat{q} rather than \hat{v}

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha(G_t^\lambda - \hat{q}(S_t, A_t, \mathbf{w}_t))\nabla\hat{q}(S_t, A_t, \mathbf{w}_t).$$

- The backward view of this algorithm is

$$\delta_t = R_{t+1} + \gamma\hat{q}(S_{t+1}, A_{t+1}, \mathbf{w}_t) - \hat{q}(S_t, A_t, \mathbf{w}_t),$$

$$\mathbf{z}_t = \gamma\lambda\mathbf{z}_{t-1} + \nabla\hat{q}(S_t, A_t, \mathbf{w}_t),$$

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \alpha\delta_t\mathbf{z}_t,$$

with $\mathbf{z}_{-1} = \mathbf{0}$.

SARSA(λ) with binary features and linear approximation

SARSA(λ) with binary features and linear approximation

Input: $\alpha > 0$, $\varepsilon > 0$, function $\mathcal{F}(s, a)$ returning active features

Output: approximate of q_* and π_*

Initialization

$\mathbf{w} \leftarrow$ arbitrarily

Loop

initialize S

$A \leftarrow \varepsilon\text{-greedy}(\hat{q}(S, \cdot, \mathbf{w}))$

for each step of the episode **do**

take action A and observe R, S'

$\delta \leftarrow R$

for $i \in \mathcal{F}(S, A)$ **do**

$\delta \leftarrow \delta - w_i$

$z_i \leftarrow z_i + 1$ or $z_i \leftarrow 1$

if S' is terminal **then**

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta \mathbf{z}$

proceed to next episode

$A' \leftarrow \varepsilon\text{-greedy}(\hat{q}(S', \cdot, \mathbf{w}))$

for $i \in \mathcal{F}(S', A')$ **do**

$\delta \leftarrow \delta + \gamma w_i$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha \delta \mathbf{z}$

$\mathbf{z} \leftarrow \gamma \lambda \mathbf{z}$

$S \leftarrow S'$

$A \leftarrow A'$

True online SARSA(λ)

True online SARSA(λ) algorithm

Input: $\alpha > 0$, $\lambda > 0$, feature function \mathbf{x} such that $\mathbf{x}(\text{terminal}, \cdot) = 0$

Output: q_* , π_*

Initialization

$\mathbf{w} \leftarrow \text{arbitrarily}$

Loop

initialize S ; $A \leftarrow \varepsilon\text{-greedy}(\hat{q}(S, \cdot, \mathbf{w}))$

$\mathbf{x} \leftarrow \mathbf{x}(S, A)$

$\mathbf{z} \leftarrow 0$

$Q_{\text{old}} \leftarrow 0$

for each step of the episode **do**

take action A and observe R, S'

$A' \leftarrow \varepsilon\text{-greedy}(\hat{q}(S', \cdot, \mathbf{w}))$

$\mathbf{x}' \leftarrow \mathbf{x}(S', A')$

$Q \leftarrow \mathbf{w}^\top \mathbf{x}$

$Q' \leftarrow \mathbf{w}^\top \mathbf{x}'$

$\delta \leftarrow R + \gamma Q' - Q$

$\mathbf{z} \leftarrow \gamma \lambda \mathbf{z} + (1 - \alpha \gamma \lambda \mathbf{z}^\top \mathbf{x}) \mathbf{x}$

$\mathbf{w} \leftarrow \mathbf{w} + \alpha (\delta + Q - Q_{\text{old}}) \mathbf{z} - \alpha (Q - Q_{\text{old}}) \mathbf{x}$

$Q_{\text{old}} \leftarrow Q'$

$\mathbf{x} \leftarrow \mathbf{x}'$

$A \leftarrow A'$

if if S' is terminal **then**

reinitialize the episode