

# Machine learning for scientific applications

Andrea Murari and Michela Gelfusa

# Overview

---


- ▶ Machine learning concepts
- ▶ Machine learning techniques
  - ▶ Classification
  - ▶ Regression
  - ▶ Advanced methods
- ▶ Applications in nuclear fusion
  - ▶ Disruption prediction in the JET tokamak
  - ▶ Causality detection



# Scientific Reliability

Ioannidis's 2005 paper "*Why Most Published Research Findings Are False*" has been the most downloaded technical paper from the journal [\*PLoS Medicine\*](#). In this paper he shows that even in the 1% of the top journals in medicine, 2/3 of the studies are contradicted by others within a few years.

Two main fundamental reasons behind this collapse of the peer review system:  
complexity of the problems and amounts of data.




# Complexity of Thermonuclear Plasmas

Difficulty to achieve robust results is an undeniable issue in the study of complex systems

The complexity of magnetised plasmas is due to:

- A different matter state far from equilibrium
- Many variables, long range interactions, high uncertainties
- Nonlinear phenomena, non separability



It is therefore very difficult to derive models from first principle. This leads to a hierarchy of descriptions of plasmas (particle, kinetic, fluid) and to a plethora of ad hoc models of limited applicability

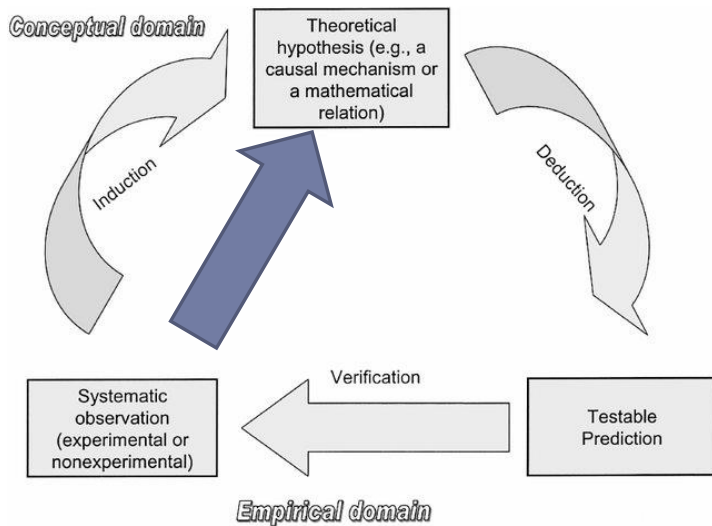
# Data Deluge and Measurements

- The amount of data produced by modern societies is enormous
- JET can produce more than 55 Gbytes of data per shot (potentially about 1 Terabyte per day). Total Warehouse: almost 0.5 Petabyte
- ATLAS can produce up to about 10 Petabytes of data per year
- Hubble Space Telescope in its prime sent to earth up to 5 Gbytes of data per day
- Commercial DVD 4.7 Gbytes (Blue Ray 50 Gbytes).

These amounts of data cannot be analysed manually in a reliable way. Given the complexity of the phenomena to be studied, there is scope for the development of new data analysis tools particularly in support to theory formulation!!

# Scientific cycle

- The “scientific process” relies on the formulation of testable predictions, which implies a dialectic relation between two domains: conceptual and empirical.



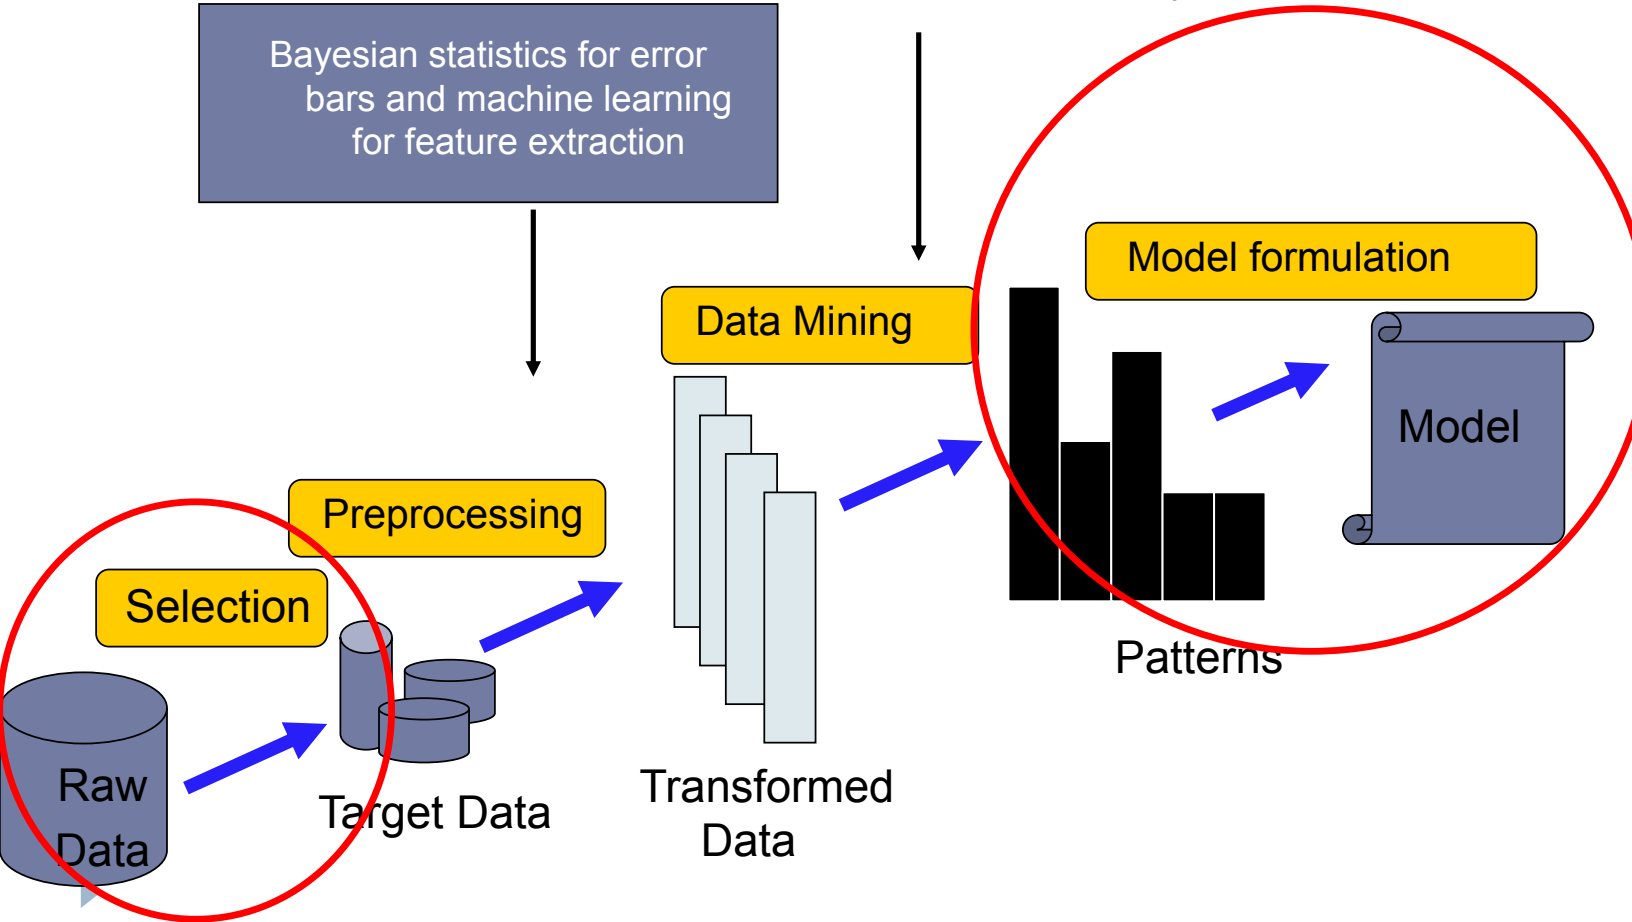
- The deduction step is very well formalised
- The induction step is more an art than a science and would benefit from: 1) more flexible tools for knowledge discovery 2) a more solid mathematization of the procedures. Data Driven Theory

1. A. Murari et al, Entropy 2017, 19, 569; doi:10.3390/e19100569
2. A. Murari et al [Nuclear Fusion, Volume 57, Number 1](#) November 2016
3. A. Murari et al 2016 Nucl. Fusion 56 076008
4. A. Murari et al 2017 Nucl. Fusion 57 126057
5. A. Murari et al [Nuclear Fusion, Volume 56, Number 2](#) (2015)
6. A. Murari et al [Plasma Physics and Controlled Fusion](#) (2015), [57](#) (1),


1. A. Murari et al 2016, Nuclear Fusion 56
2. A. Murari et al. Nuclear. Fusion **57** (2017) 016024 2017,
3. A. Murari et al. Nuclear Fusion , **57**, Number 12, September 2017
4. A. Murari et al [Nuclear Fusion, Volume 58, Number 5](#), March 2018

# Data Analysis: an overview

Given the complexity of the inference process, various machine learning tools are used to handle large amounts of data, which is often unstructured.



# Machine learning: concepts



- ▶ Learning does not mean '*learning by heart*' (any computer can memorize)
- ▶ Learning means '*generalization capability*': we learn with some samples and predict for other samples







# Machine learning: concepts

---

- ▶ The learning problem is the problem of finding a desired dependence (function) using a *limited* number of observations (training data)
  - ▶ Classification: the function represents the separation frontier between two classes
  - ▶ Regression: the function provides a fit to the data



Classification

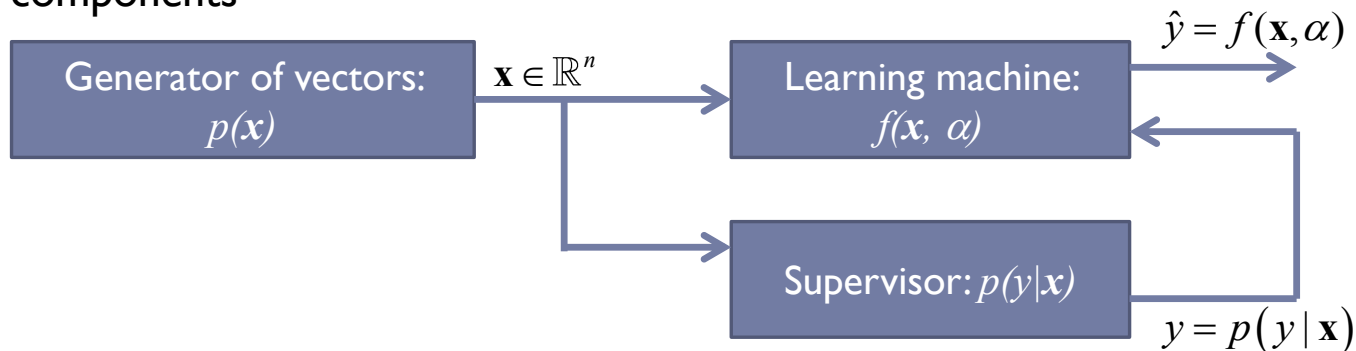


Regression



# Machine learning: concepts

- ▶ The general model of *learning from examples* is described through three components



$p(\mathbf{x})$ : fixed but unknown probability distribution function

$y = p(y|\mathbf{x})$  (fixed and unknown)


$f(\mathbf{x}, \alpha)$ :  $\alpha$  indicates an index in the class of functions considered

$(\mathbf{x}_i, y_i), i = 1, \dots, N$ : training samples

- ▶ The problem of learning is that of choosing from the given set of functions  $f(\mathbf{x}, \alpha)$ , the one that best approximates the supervisor's response


$$\{\hat{y}_1, \hat{y}_2, \dots, \hat{y}_N\} \text{ "close" to } \{y_1, y_2, \dots, y_N\}$$

# Classification example




$$f(x, \alpha) = \begin{cases} 0 & \text{if } x_2 < ax_1 + b \\ 1 & \text{if } x_2 \geq ax_1 + b \end{cases}$$

$$\alpha = (a, b)$$



# Regression example



$$f(x, \alpha) = \phi_1(x, A_1, f_1), \quad \alpha = (A_1, f_1)$$

$$f(x, \beta) = \phi_2(x, A_2, f_2), \quad \beta = (A_2, f_2)$$

$$f(x, \gamma) = \phi_3(x, m, b), \quad \gamma = (m, b)$$



---

# Classification



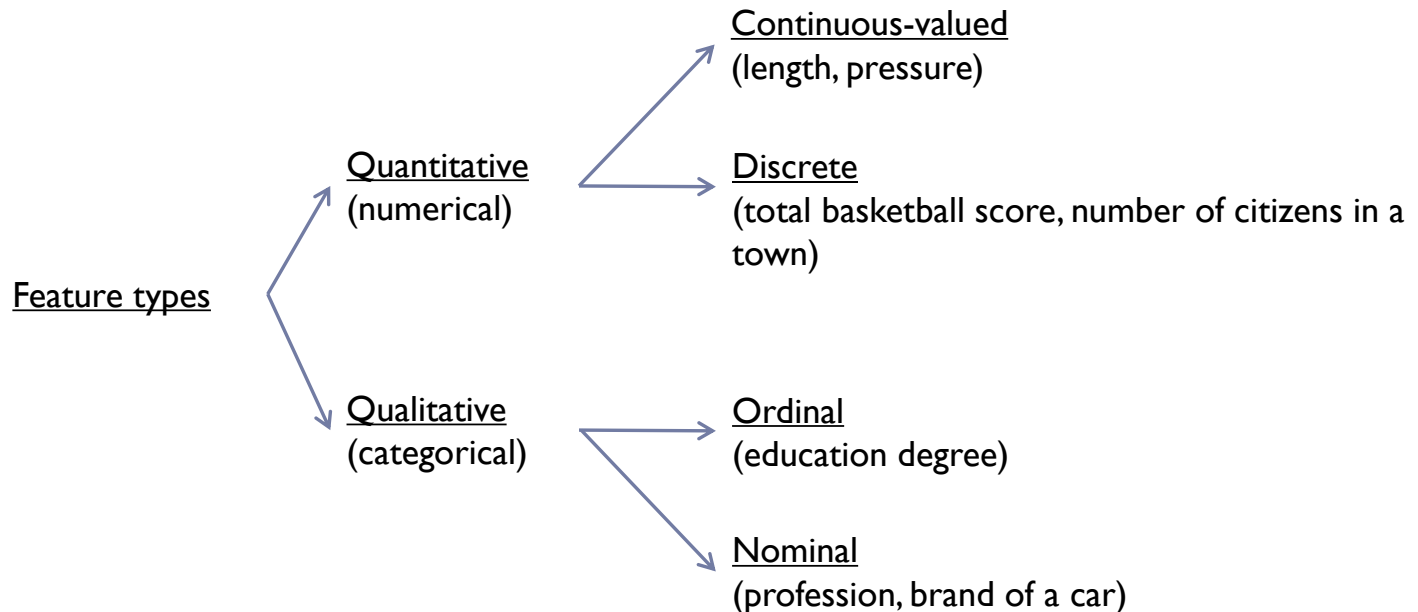
# Description of objects

---

**Dataset:**  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_i, y_i), \dots, (\mathbf{x}_N, y_N)$

$\mathbf{x}_i \in \mathbb{R}^m$ : features that are of distinctive nature (object description with attributes managed by computers)

$y_i \in \{L_1, L_2, \dots, L_K\}$ : label of the sample  $\mathbf{x}_i$



# Description of objects

## ► Examples of feature vectors

### ► Disruptions

$$\mathbf{x}_1 = (\beta_p(t_s), I_p(t_s), n_e(t_s), \dots), \quad y_1 \in \{D, N\}$$

$$\mathbf{x}_2 = (\beta_p(t_s + T), I_p(t_s + T), n_e(t_s + T), \dots), \quad y_2 \in \{D, N\}$$

$$\mathbf{x}_3 = (\beta_p(t_s + 2T), I_p(t_s + 2T), n_e(t_s + 2T), \dots), \quad y_3 \in \{D, N\}$$

...



### ► L/H transition

$$(\mathbf{x}_i, y_i), \quad \mathbf{x}_i \in \mathbb{R}^m, \quad y_i \in \{L, H\}$$

### ► Image classification

- $\mathbf{x}_i$ : the set of pixels of an image

- $y_i \in \{1, 2, 3\}$





# Supervised classifiers

Dataset:  $(x_1, y_1), (x_2, y_2), \dots, (x_i, y_i), \dots, (x_N, y_N)$

$x_i \in \mathbf{R}^m$ : features that are of distinctive nature (object description with attributes managed by computers)  
 $y_i \in \{L_1, L_2, \dots, L_K\}$ : known label of the sample  $x_i$

**Objective:** to determine a separating function between classes (generalization) to predict the label of new samples with known feature vectors  $((x_{N+1}, y_{N+1}), (x_{N+2}, y_{N+2}), \dots)$





# Supervised classifiers

---

- ▶ **Single classifiers**
  - ▶ Neural networks
  - ▶ Support Vector Machines (SVM)
  - ▶ Bayes decision theory
    - ▶ Parametric method
    - ▶ Non-parametric method
  - ▶ Classification trees
- ▶ **Combining classifiers**



# Introduction to neural networks

---

- ▶ **What is an (artificial) neural network**
  - ▶ A large set of **nodes** (units, neurons, processing elements)
    - ▶ Each node has input and output
    - ▶ Each node performs a **simple** computation by its node function
  - ▶ **Weighted connections** between nodes
    - ▶ Connectivity gives the structure/architecture of the net
    - ▶ Connections/links have directions
    - ▶ What can be computed by a NN is primarily determined by the connections and their weights
  - ▶ A very much simplified version of networks of neurons in animal nerve systems



# Introduction

---

## **Von Neumann machine**

---

- ▶ One or a few high speed (nm second) processors with considerable computing power
- ▶ One or a few shared high speed buses for communication
- ▶ Sequential memory access by address
- ▶ Problem-solving knowledge is separated from the computing component
- ▶ Hard to be adaptive


## **Human Brain**

---

- Large # ( $10^{11}$ ) of low speed processors (ms) with limited computing power
- Large # ( $10^{15}$ ) of low speed connections
- Content addressable recall (CAM)
- Problem-solving knowledge resides in the connectivity of neurons
- Adaptation by changing the connectivity



## • Biological neural activity



- Each neuron has a *body*, an *axon*, and many *dendrites*
  - Can be in one of the two states: *firing* and *rest*.
  - Neuron fires if the total incoming stimulus exceeds the threshold
- *Synapse*: thin gap between axon of one neuron and dendrite of another.
  - Signal exchange
  - Synaptic strength/efficiency



# Introduction

## ANN

- ▶ **Nodes**
  - ▶ input
  - ▶ output
  - ▶ node function
- ▶ **Connections**
  - ▶ connection strength

## Bio NN


- **Cell body**
  - signals from other neurons
  - firing frequency
  - firing mechanism
- **Synapses**
  - synaptic strength

- Highly parallel, simple local computation (*at neuron level*) achieves global results as emerging property of the interaction (*at network level*)
- Pattern directed (meaning of individual nodes only in the context of a pattern)
- Fault-tolerant/graceful degrading
- Learning/adaptation plays important role.

# ANN Neuron Models



- ▶ Node Input
  - ▶ Each node has one or more inputs from other nodes, and one output to other nodes
  - ▶ Input/output values can be
    - ▶ Binary  $\{0, 1\}$ ; Bipolar  $\{-1, 1\}$ ; or Continuous (bounded or not)
- ▶ Weighted sum of inputs

$$\text{output} = f(\text{net}) \text{ where } \text{net} = \sum_{i=1}^n w_i x_i$$



# Node Functions

- Node functions (linear)
  - Identity function:  $f(net) = net$ .
  - Constant function:  $f(net) = c$ .
- Node functions (**non-linear**)
  - Step functions and ramp functions
  - Sigmoid functions (**differentiable**)




## ► Sigmoid function

- S-shaped
- Continuous and everywhere differentiable
- Rotationally symmetric
- Asymptotically approaches saturation points
- Examples:
  - Logistic function

$$o_j = \frac{1}{1 + e^{-net_j}}$$

- Hyperbolic tangent

$$o_j = \frac{e^{net_j} - e^{-net_j}}{e^{net_j} + e^{-net_j}}$$






# Network Architecture

## ▶ (Asymmetric) Fully Connected Networks

- ▶ Every node is connected to every other node
- ▶ Connection may be excitatory (positive), inhibitory (negative), or irrelevant ( $\approx 0$ ).
- ▶ Most general
- ▶ Symmetric fully connected nets: weights are symmetric ( $w_{ij} = w_{ji}$ )



**Input nodes:** receive input from the environment


**Output nodes:** send signals to the environment

**Hidden nodes:** no direct interaction to the environment

# Network Architecture

## ► Layered Networks

- Nodes are partitioned into subsets, called layers.
- No connections that lead from nodes in layer  $j$  to those in layer  $k$  if  $j > k$ .




- Inputs from the environment are applied to nodes in layer 0 (**input layer**).
- Nodes in input layer are place holders with no computation occurring (i.e., their node functions are identity function)

# Network Architecture

## ► Feedforward Networks

- A connection is allowed from a node in layer  $i$  only to nodes in layer  $i + 1$ .
- Most widely used architecture.




Conceptually, nodes at higher levels successively abstract features from preceding layers

# Supervised classifiers: artificial neural networks

Samples:  $(\mathbf{x}_j, y_j), \mathbf{x}_j \in \mathbb{R}^n, j = 1, \dots, N, y_j \in \{L_1, L_2\}$

Aim: determine  $\mathbf{W}$  for a minimum error



$$s(u) = \frac{1}{1 + \exp(-u)}$$

$$s(u) = \tanh(u) = \frac{2}{1 + \exp(-2u)} - 1$$

Input

Hidden layer 1

Hidden layer 2

Output

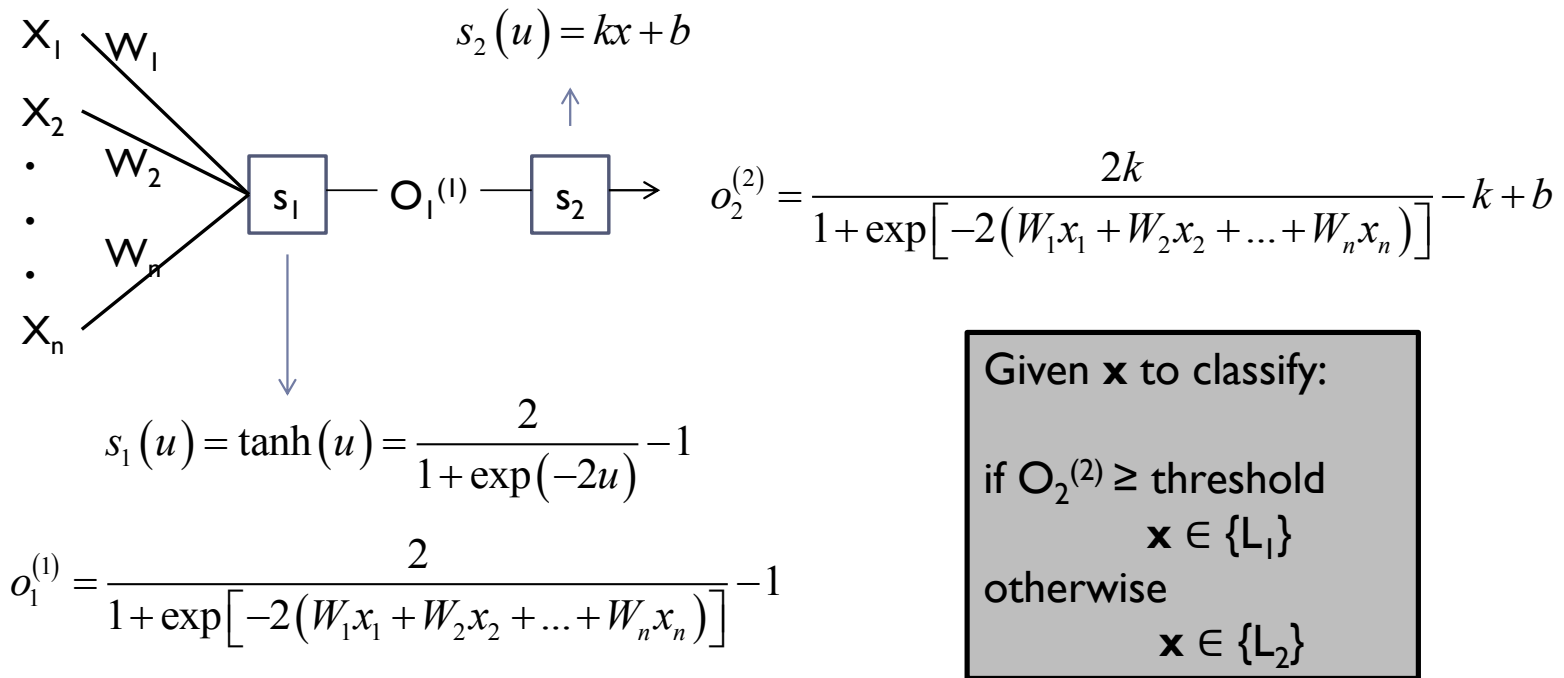
$\mathbf{W}^{(1)}$ :  $(n+1 \times k)$  matrix

$\mathbf{W}^{(2)}$ :  $(k+1 \times m)$  matrix

$\mathbf{W}^{(3)}$ :  $(m+1 \times 1)$  matrix

# Supervised classifiers: artificial neural networks (example)

Samples:  $(\mathbf{x}_j, y_j), \mathbf{x}_j \in \mathbb{R}^n, j = 1, \dots, N, y_j \in \{L_1, L_2\}$



# Neural Network Learning

---

- ▶ **Real neural learning**

- ▶ Synapses change size and strength with experience.
- ▶ **Hebbian learning**: When two connected neurons are firing at the same time, the strength of the synapse between them increases.

“Neurons that fire together, wire together.”

- ▶ **ANN learning**

- ▶ Construct the NN by training using samples
- ▶ Training is primarily to change the weights
- ▶ We will look at two particular training methods
  - ▶ **Perceptron**
  - ▶ **backpropagation**





# Perceptron

- ▶ Structure: single layer
- ▶ Non-linear output nodes

Threshold units

$$o_j = \begin{cases} 0 & \text{if } net_j < T_j \\ 1 & \text{if } net_j \geq T_j \end{cases}$$

- ▶ Supervised learning
  - ▶ Learning weights so that output node produces correct output for each training sample



# Perceptron Learning Rule

---

- ▶ Update weights by:

$$w_{ji} = w_{ji} + \eta(t_j - o_j)o_i$$

where  $\eta$  is the “learning rate”

$t_j$  is the teacher specified output for unit  $j$

$(t_j - o_j)$  is the **error** (training is error driven)

- ▶ Equivalent to rules:

- ▶ If output is correct do nothing.
- ▶ If output is high, lower weights on active inputs
- ▶ If output is low, increase weights on active inputs





# Perceptron Learning Algorithm

---

- ▶ Iteratively update weights until convergence.

Initialize weights to random values

Until outputs of all training examples are correct

For each training pair,  $E$ , do:

    Compute current output  $o_j$  for  $E$  given its inputs

    Compare current output to target value,  $t_j$ , for  $E$


    Update synaptic weights using learning rule

- ▶ Each execution of the outer loop is typically called an *epoch*.



# Perceptron as a Linear Separator

- ▶ Since perceptron uses linear threshold function, it is searching for a linear separator that discriminates the classes.




$$w_{12}o_2 + w_{13}o_3 > T_1$$

$$o_3 > -\frac{w_{12}}{w_{13}}o_2 + \frac{T_1}{w_{13}}$$

**Or hyperplane in  
n-dimensional space**

# Concept Perceptron Cannot Learn

- ▶ Cannot learn exclusive-or, or parity function in general because they are not **linearly separable**.



# Perceptron Convergence Theorem

---

- ▶ **Perceptron convergence theorem:** If the data is linearly separable and therefore a set of weights exist that are consistent with the data, then the Perceptron algorithm will eventually converge to a consistent set of weights.
- ▶ **Perceptron cycling theorem:** If the data is not linearly separable, the Perceptron algorithm will eventually repeat a set of weights and threshold at the end of some epoch and therefore enter an infinite loop.
  - ▶ By checking for repeated weights+threshold, one can guarantee termination with either a positive or negative result.



# Perceptron Limits


---

- ▶ System obviously cannot learn concepts it cannot represent.
- ▶ Minsky and Papert (1969) wrote a book analyzing the perceptron and demonstrating many functions it could not learn.
- ▶ These results discouraged further research on neural nets; and symbolic AI became the dominate paradigm.



# Multi-Layer Feed-Forward Networks

- ▶ A typical multi-layer network consists of an input layer, one or more hidden and output layer, each fully connected to the next, with activation feeding forward.




- ▶ **Nodes at hidden layers MUST be non-linear** (typically a sigmoid function)
- ▶ Given an arbitrary number of hidden units with a single hidden layer, there exists a set of weights which can compute any given Boolean function (or any L2 function). **But an effective learning algorithm for such networks was thought to be difficult.**



# Sample Learned XOR Network

---



Hidden Unit A represents:  $\neg(X \wedge Y)$

Hidden Unit B represents:  $\neg(X \vee Y)$

Output O represents:  $A \wedge \neg B = \neg(X \wedge Y) \wedge (X \vee Y)$   
 $= X \oplus Y$



# Gradient Descent

---

- ▶ Define objective to minimize error:

$$E(W) = \sum_{d \in D} \sum_{k \in K} (t_{kd} - o_{kd})^2$$

where  $D$  is the set of training examples,  $K$  is the set of output units,  $t_{kd}$  and  $o_{kd}$  are, respectively, the teacher and current output for unit  $k$  for example  $d$ .

- ▶ The derivative of a sigmoid unit with respect to net input is:

$$\frac{\partial o_j}{\partial net_j} = o_j(1 - o_j)$$

- ▶ Learning rule to change weights to minimize error is:

$$\Delta w_{ji} = -\eta \frac{\partial E}{\partial w_{ji}}$$





# Backpropagation Learning Rule

---

- ▶ Each weight changed by:

$$\Delta w_{ji} = \eta \delta_j o_i$$

$$\delta_j = o_j(1 - o_j)(t_j - o_j) \quad \text{if } j \text{ is an output unit}$$

$$\delta_j = o_j(1 - o_j) \sum_k \delta_k w_{kj} \quad \text{if } j \text{ is a hidden unit}$$

where  $\eta$  is a constant called the learning rate

$t_j$  is the correct teacher output for unit  $j$

$\delta_j$  is the error measure for unit  $j$



# Error Backpropagation

---

- ▶ First calculate error of output units and use this to change the top layer of weights.

Current output:  $o_j=0.2$


Correct output:  $t_j=1.0$

Error  $\delta_j = o_j(1-o_j)(t_j-o_j)$

$0.2(1-0.2)(1-0.2)=0.128$

Update weights into  $j$


$$\Delta w_{ji} = \eta \delta_j o_i$$



# Error Backpropagation

---


- ▶ Next calculate error for hidden units based on errors on the output units it feeds into.



# Error Backpropagation

---

- ▶ Finally update bottom layer of weights based on errors calculated for hidden units.



# Backpropagation Training Algorithm

---

Create the 3-layer network with  $H$  hidden units with full connectivity between layers. Set weights to small random real values.

Until all training examples produce the correct value (within  $\varepsilon$ ), or mean squared error ceases to decrease, or other termination criteria:

- Begin epoch

- For each training example,  $d$ , do:

  - Calculate network output for  $d$ 's input values

  - Compute error between current output and correct output for  $d$

    - Update weights by backpropagating error and using learning rule

- End epoch



# Hidden Unit Representations

---

- ▶ Trained hidden units can be seen as newly constructed features that make the target concept linearly separable in the transformed space.
- ▶ On many real domains, hidden units can be interpreted as representing meaningful features such as vowel detectors or edge detectors, etc..
- ▶ However, the hidden layer can also become a distributed representation of the input in which each individual unit is not easily interpretable as a meaningful feature.



# Representational Power

---

- ▶ **Boolean functions:** Any boolean function can be represented by a two-layer network with sufficient hidden units.
- ▶ **Continuous functions:** Any bounded continuous function can be approximated with arbitrarily small error by a two-layer network.
  - ▶ Sigmoid functions can act as a set of basis functions for composing more complex functions, like sine waves in Fourier analysis.
- ▶ **Arbitrary function:** Any function can be approximated to arbitrary accuracy by a three-layer network.



# Successful Applications

---

- ▶ Text to Speech (NetTalk)
- ▶ Fraud detection
- ▶ Financial Applications
  - ▶ HNC (eventually bought by Fair Isaac)
- ▶ Chemical Plant Control
  - ▶ Pavillion Technologies
- ▶ Automated Vehicles
- ▶ Game Playing
  - ▶ Neurogammon
- ▶ Handwriting recognition





# Hand-written character recognition

- MNIST: a data set of hand-written digits
  - 60,000 training samples
  - 10,000 test samples
  - Each sample consists of  $28 \times 28 = 784$  pixels

	Failure rate for test samples
• Various techniques have been tried	
– Linear classifier:	12.0%
– 2-layer BP net (300 hidden nodes)	4.7%
– 3-layer BP net (300+200 hidden nodes)	3.05%
– Support vector machine (SVM)	1.4%
– Convolutional net	0.4%
– <b>6 layer BP net (7500 hidden nodes):</b>	<b>0.35%</b>

# Strengths of BP Learning

---

- ▶ **Great representation power**

- ▶ Any meaningful function can be represented by a BP net
- ▶ Many such functions can be approximated by BP learning (gradient descent approach)

- ▶ **Easy to apply**

- ▶ Only requires a good set of training samples
- ▶ Does not require substantial prior knowledge or deep understanding of the domain itself (ill structured problems)
- ▶ Tolerates noise and missing data in training samples (graceful degrading)

- ▶ **Easy to implement** the core of the learning algorithm

- ▶ **Good generalization power**

- ▶ Often produces accurate results for inputs outside the training set



# Deficiencies of BP Learning

---

- ▶ Learning often takes a **long time** to converge
  - ▶ Complex functions often need hundreds or thousands of epochs
- ▶ The network is essentially a **black box**
  - ▶ It may provide a desired mapping between input and output vectors ( $\mathbf{x}$ ,  $\mathbf{o}$ ) but does not have the information of why a particular  $\mathbf{x}$  is mapped to a particular  $\mathbf{o}$ .
  - ▶ It thus cannot provide an intuitive (e.g., causal) explanation for the computed result.
  - ▶ This is because the hidden nodes and the learned weights do not have clear semantics.
    - ▶ What can be learned are operational parameters, not general, abstract knowledge of a domain
  - ▶ Unlike many statistical methods, there is no theoretically well-founded way to **assess the quality** of BP learning
    - ▶ What is the confidence level of  $\mathbf{o}$  computed from input  $\mathbf{x}$  using such net?
    - ▶ What is the confidence level for a trained BP net, with the final training/test error (which may or may not be close to zero)?

# Deficiencies of BP Learning

---

- ▶ Problem with gradient descent approach
  - ▶ only guarantees to reduce the total error to a **local minimum**. ( $E$  might not be reduced to zero)
    - ▶ Cannot escape from the local minimum error state
    - ▶ **Not every function that is representable can be learned**
  - ▶ How bad: depends on the shape of the error surface. Too many valleys/wells will make it easy to be trapped in local minima
  - ▶ Possible remedies:
    - ▶ Try nets with different # of hidden layers and hidden nodes (they may lead to different error surfaces, some might be better than others)
    - ▶ Try different initial weights (different starting points on the surface)
    - ▶ Forced escape from local minima by random perturbation (e.g., **simulated annealing**)

- ▶ **Generalization** is not guaranteed even if the error is reduced to 0

- ▶ **Over-fitting**/over-training problem: trained net fits the training samples perfectly ( $E$  reduced to 0) but it does not give accurate outputs for inputs not in the training set

– Possible remedies:

- More and better samples
- Using smaller net if possible
- Using larger error bound (forced early termination)
- Introducing noise into samples
  - modify  $(x_1, \dots, x_n)$  to  $(x_1 + \alpha_1, \dots, x_n + \alpha_n)$  where  $\alpha_i$  are small random displacements

- **Cross-Validation**

- leave some (~20%) samples as test data (not used for weight update)
- periodically check error on test data
- learning stops when error on test data starts to increase

